

Question 1

P1 Tree Model -- Step 1: Load Data

Setup the environment to load the tree and randomForest libs and load uscrime.txt into data_df:

```
# Clear the environment
rm(list = ls())

# Comment in set.seed(33) to repeat results
set.seed(33)

# Load tree lib
require(tree)
require(randomForest)

# Load crime data into a data frame
data_df <- read.table("uscrime.txt", header=TRUE)
```

P1 Tree Model -- Step 2: Train Tree Model

Using the tree function, I trained the tree_model, visualized it, and calculated the R²:

```
# Train tree model
tree_model <- tree(Crime ~., data_df)
summary(tree_model)
```

Regression tree:

```
tree(formula = Crime ~ ., data = data_df)
```

Variables actually used in tree construction:

```
[1] "Po1" "Pop" "LF" "NW"
```

Number of terminal nodes: 7

Residual mean deviance: 47390 = 1896000 / 40


Distribution of residuals:

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-573.900	-98.300	-1.545	0.000	110.600	490.100

```
# Visualize tree model
```

```
plot(tree_model)
```

```
text(tree_model)
```



```

# Function to calculate R^2
ComputerR2 <- function(yhat_df, data_df) {
  SSres <- sum((yhat_df - data_df$Crime)^2)
  SStot <- sum((data_df$Crime - mean(data_df$Crime))^2)
  R2 <- 1 - SSres/SStot
  return(R2)
}

# Calculate R^2
tree_yhat <- predict(tree_model)
tree_r2 <- ComputerR2(tree_yhat, data_df)
Tree_r2 # 0.7244962

```

P1 Tree Model -- Step 4: Prune the Tree

Due to the limited dataset size, the tree should be pruned to a smaller size to allow enough data at each terminal leaf:

```
# Manually prune tree
```

```
tree_model_pruned <- prune.tree(tree_model,best = 4)  
summary(tree_model_pruned)
```

Regression tree:

```
snip.tree(tree = tree_model, nodes = c(6L, 2L))
```

Variables actually used in tree construction:

```
[1] "Po1" "NW"
```

Number of terminal nodes: 4

Residual mean deviance: 61220 = 2633000 / 43


Distribution of residuals:

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-573.90	-152.60	35.39	0.00	158.90	490.10

```
# Visualize pruned tree
```

```
plot(tree_model_pruned)
```

```
text(tree_model_pruned)
```



```
# Calc R^2
pruned_yhat <- predict(tree_model_pruned)
pruned_r2 <- ComputerR2(pruned_yhat, data_df)
Pruned_r2 # 0.6174017
```

P2 Forest Model -- Step 1: Train Forest Model

First I created a variable to set my factor set to $1 + \log(n)$ and then trained a forest model using `randomForest()`:

```
# Use 1+log(n) standard to pick number of factors in each set
factor_set <- round(1 + log(ncol(data_df)))

# Train forest model
forest_model <- randomForest(Crime ~., data_df, mtry = factor_set,
importance = TRUE, ntree = 500)
Forest_model
```

Call:

```
randomForest(formula = Crime ~ ., data = data_df, mtry = factor_set,
importance = TRUE, ntree = 500)
      Type of random forest: regression
      Number of trees: 500
No. of variables tried at each split: 4
```

```
      Mean of squared residuals: 83777.12
      % Var explained: 42.78
```

P2 Forest Model -- Step 2: Assess Forest Model

I then calculated R^2 and visualized variable importance for the forest model:

```
# Use R2 function to calculate
forest_model_yhat <- predict(forest_model)
ComputerR2(forest_model_yhat, data_df)
# 0.421717
```

```
# Visualize variable importance
importance(forest_model)
varImpPlot(forest_model)
```

	%IncMSE	IncNodePurity
M	2.0028307	226264.31

So	2.5749953	29100.98
Ed	3.0818775	302023.75
Po1	11.5680884	1198828.12
Po2	11.7805096	1015636.06
LF	5.0612871	283691.24
M.F	1.2939713	257229.86
Pop	0.6063912	368227.18
NW	8.3066286	472483.85
U1	0.7838364	131504.61
U2	2.2275446	196529.76
Wealth	4.0709227	611259.13
Ineq	1.2877040	231915.88
Prob	7.7684823	734271.60
Time	1.8736891	195756.99

Step 4: Interpret Results

The `tree_model_pruned` is limited in its predictive ability because it assigns one of four predictions; limiting its variability to new data. Whereas, the `forest_model` has a lot more variability in its predicted values but is harder to explain. The `tree_model_pruned` does have a higher R^2 but that's likely due to overfitting on the small `uscrime` dataset.

Use models to predict

```
predict(tree_model_pruned, data_df[,1:15])
predict(forest_model, data_df[,1:15])
```

```
> predict(tree_model_pruned, data_df[,1:15])
 1      2      3      4      5      6      7      8      9     10     11     12     13     14     15     16
669.6087 1502.8750 669.6087 1502.8750 886.9000 886.9000 1041.0000 1502.8750 669.6087 669.6087 1502.8750 669.6087 669.6087 669.6087 1041.0000 669.6087 1041.0000
17      18      19      20      21      22      23      24      25      26      27      28      29      30      31      32
669.6087 1502.8750 886.9000 1502.8750 669.6087 669.6087 1041.0000 886.9000 669.6087 1502.8750 669.6087 1041.0000 1502.8750 669.6087 669.6087 1041.0000
33      34      35      36      37      38      39      40      41      42      43      44      45      46      47
669.6087 886.9000 886.9000 886.9000 669.6087 669.6087 669.6087 1041.0000 669.6087 669.6087 669.6087 886.9000 669.6087 886.9000 886.9000
> predict(forest_model, data_df[,1:15])
 1      2      3      4      5      6      7      8      9     10     11     12     13     14     15     16
783.7498 1384.1008 597.4483 1664.2102 1150.0543 943.2849 948.0244 1329.0192 826.6916 719.7465 1534.9990 841.8490 582.8568 670.5281 770.4288 958.9704
17      18      19      20      21      22      23      24      25      26      27      28      29      30      31      32
596.1851 986.2564 889.1561 1194.8476 765.3563 559.6587 1149.3142 947.6447 591.4752 1636.3687 580.2940 1162.6364 1183.8141 757.2752 534.6020 889.0966
33      34      35      36      37      38      39      40      41      42      43      44      45      46      47
929.0435 939.9612 860.4663 1204.9387 805.1989 581.4471 815.8237 1151.5587 850.5561 548.8967 837.1231 1039.4939 526.9124 768.5093 918.2002
```

Question 2

A situation where a logistic regression model could be useful in my personal life is whether I will enjoy seeing a movie in the theater or not. Potential predictors for this model include: weeks since release, box office receipts, Rotten Tomatoes score, IMDB score, and genre.

Question 3.1

Step 1: Load Data

First, I cleared the environment, set the seed, and loaded germancredit.txt into data_df. Once loaded, the response (V21) needed to be updated to 0 & 1 instead of 1 & 2. I also split the dataset into a training and testing set:

```
# Clear the environment
rm(list = ls())

# Comment in set.seed(33) to repeat results
set.seed(33)

# Load crime data into a data frame
data_df <- read.table("germancredit.txt", header=FALSE)

# Change Response (V21) to 1 (good) or 0 (bad)
data_df$V21[data_df$V21 == 1] <- 1
data_df$V21[data_df$V21 == 2] <- 0

# Create training and test datasets
pct70 <- sample(1:nrow(data_df), size = round(0.7*(nrow(data_df))))
train_df <- data_df[pct70,]
test_df <- data_df[-pct70,]
```

Step 2: Build Model with All Factors

I started off by building the model using all factors. This helped to determine which factors were insignificant and could be removed:

```
# Build model with all factors to determine significant factors
model_all <- glm(V21 ~., family = binomial(link="logit"), train_df)
summary(model_all)
Call:
glm(formula = V21 ~ ., family = binomial(link = "logit"), data = train_df)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-2.6666  -0.6807   0.3539   0.6914   2.3340
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)	
(Intercept)	-1.238e+00	1.322e+00	-0.937	0.348800	
V1A12	3.811e-01	2.608e-01	1.461	0.143946	
V1A13	9.413e-01	4.369e-01	2.155	0.031189	*
V1A14	1.799e+00	2.880e-01	6.248	4.15e-10	***
V2	-2.844e-02	1.075e-02	-2.645	0.008178	**
V3A31	6.488e-02	6.742e-01	0.096	0.923331	
V3A32	9.223e-01	5.151e-01	1.791	0.073351	.
V3A33	1.409e+00	5.737e-01	2.455	0.014082	*
V3A34	1.672e+00	5.269e-01	3.173	0.001508	**
V4A41	1.729e+00	4.716e-01	3.666	0.000246	***
V4A410	1.244e+00	8.354e-01	1.489	0.136548	
V4A42	7.339e-01	3.141e-01	2.337	0.019451	*
V4A43	1.133e+00	3.031e-01	3.738	0.000185	***
V4A44	5.839e-01	8.025e-01	0.728	0.466866	
V4A45	9.406e-01	7.074e-01	1.330	0.183606	
V4A46	-7.489e-02	4.803e-01	-0.156	0.876102	
V4A48	2.142e+00	1.298e+00	1.650	0.098878	.
V4A49	9.885e-01	4.028e-01	2.454	0.014128	*
V5	-1.248e-04	5.418e-05	-2.304	0.021221	*
V6A62	5.615e-01	3.540e-01	1.586	0.112743	
V6A63	1.135e+00	5.757e-01	1.971	0.048695	*
V6A64	1.122e+00	6.142e-01	1.827	0.067680	.
V6A65	1.053e+00	3.263e-01	3.226	0.001256	**
V7A72	2.824e-01	5.257e-01	0.537	0.591063	
V7A73	1.380e-01	5.093e-01	0.271	0.786451	
V7A74	9.270e-01	5.561e-01	1.667	0.095528	.
V7A75	2.524e-01	5.091e-01	0.496	0.620012	
V8	-4.225e-01	1.089e-01	-3.878	0.000105	***
V9A92	7.668e-01	4.866e-01	1.576	0.115064	
V9A93	1.125e+00	4.761e-01	2.363	0.018109	*
V9A94	7.472e-01	5.660e-01	1.320	0.186832	
V10A102	-1.883e-02	4.945e-01	-0.038	0.969619	
V10A103	9.826e-01	5.303e-01	1.853	0.063872	.
V11	-8.698e-02	1.060e-01	-0.820	0.412021	
V12A122	-4.707e-01	3.081e-01	-1.527	0.126654	
V12A123	-3.021e-01	2.860e-01	-1.056	0.290809	
V12A124	-8.060e-01	5.169e-01	-1.559	0.118977	
V13	1.909e-02	1.114e-02	1.714	0.086466	.
V14A142	2.648e-01	5.122e-01	0.517	0.605189	
V14A143	3.556e-01	2.861e-01	1.243	0.213886	

```

V15A152      6.519e-01  2.876e-01  2.267 0.023406 *
V15A153      1.105e+00  5.904e-01  1.872 0.061223 .
V16          -1.937e-01  2.371e-01  -0.817 0.413772
V17A172      -4.302e-01  8.056e-01  -0.534 0.593374
V17A173      -2.375e-01  7.761e-01  -0.306 0.759625
V17A174      -3.778e-01  7.785e-01  -0.485 0.627526
V18          -2.013e-01  3.085e-01  -0.653 0.514023
V19A192       3.302e-01  2.499e-01  1.321 0.186416
V20A202       1.678e+00  7.646e-01  2.194 0.028215 *

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 861.88 on 699 degrees of freedom
Residual deviance: 617.04 on 651 degrees of freedom
AIC: 715.04

Number of Fisher Scoring iterations: 5

Step 3: Train Refined Model

I used step() to assist with variable selection. After determining the optimal factor set, I retrained the model:

```
# Use step for variable selection
```

```
step_output <- step(model_all)
```

```
step_output
```

```
Call: glm(formula = V21 ~ V1 + V2 + V3 + V4 + V5 + V6 + V8 + V9 + V15 +
  V20, family = binomial(link = "logit"), data = train_df)
```

Coefficients:

```

(Intercept)      V1A12      V1A13      V1A14      V2
V3A31      V3A32      V3A33      V3A34      V4A41      V4A410
V4A42      V4A43
-1.3460003    0.4718098    1.0697297    1.8212661    -0.0279530
0.2423034    1.1363883    1.4692351    1.7854395    1.6885645    1.1265416
0.6139784    1.1482593
      V4A44      V4A45      V4A46      V4A48      V4A49
V5      V6A62      V6A63      V6A64      V6A65      V8
V9A92      V9A93
 0.5442421    0.8686075   -0.2715785    2.0768567    0.9240859

```


-0.0001192	0.3949389	1.0925827	1.1476477	1.0200351	-0.3802927
0.6369140	1.0034172				
V9A94	V15A152	V15A153	V20A202		
0.6844993	0.6639122	0.6466581	1.5887962		

Degrees of Freedom: 699 Total (i.e. Null); 670 Residual
Null Deviance: 861.9
Residual Deviance: 638.3 AIC: 698.3

Rebuild model with optimal step_output variable combination

```
model_refined <- glm(V21 ~ V1 + V2 + V3 + V4 + V5 + V6 + V7 + V8 + V9 + V10
+ V14 + V15 + V20, family = binomial(link="logit"), data_df)
summary(model_refined)
```

Call:

```
glm(formula = V21 ~ V1 + V2 + V3 + V4 + V5 + V6 + V7 + V8 + V9 +
V10 + V14 + V15 + V20, family = binomial(link = "logit"),
data = data_df)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-2.6294	-0.7269	0.3939	0.6910	2.3155

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-1.145e+00	7.478e-01	-1.531	0.125670
V1A12	4.210e-01	2.137e-01	1.970	0.048833 *
V1A13	1.069e+00	3.621e-01	2.952	0.003156 **
V1A14	1.751e+00	2.294e-01	7.635	2.27e-14 ***
V2	-3.038e-02	9.004e-03	-3.374	0.000740 ***
V3A31	2.944e-02	5.270e-01	0.056	0.955457
V3A32	7.672e-01	4.111e-01	1.866	0.062001 .
V3A33	8.984e-01	4.672e-01	1.923	0.054515 .
V3A34	1.447e+00	4.337e-01	3.337	0.000847 ***
V4A41	1.635e+00	3.650e-01	4.478	7.53e-06 ***
V4A410	1.623e+00	7.572e-01	2.144	0.032052 *
V4A42	7.281e-01	2.530e-01	2.878	0.004000 **
V4A43	8.826e-01	2.433e-01	3.627	0.000286 ***
V4A44	5.943e-01	7.552e-01	0.787	0.431336
V4A45	1.757e-01	5.443e-01	0.323	0.746795
V4A46	-9.063e-02	3.903e-01	-0.232	0.816397
V4A48	1.978e+00	1.219e+00	1.623	0.104535

```

V4A49      7.809e-01  3.291e-01  2.373 0.017660 *
V5         -1.156e-04 4.149e-05 -2.786 0.005340 **
V6A62      2.340e-01 2.776e-01  0.843 0.399186
V6A63      4.235e-01 3.966e-01  1.068 0.285639
V6A64      1.318e+00 5.118e-01  2.576 0.010005 *
V6A65      9.527e-01 2.577e-01  3.698 0.000218 ***
V7A72     -2.114e-01 3.729e-01 -0.567 0.570717
V7A73     -4.642e-02 3.480e-01 -0.133 0.893873
V7A74      5.869e-01 3.905e-01  1.503 0.132778
V7A75      1.118e-01 3.625e-01  0.308 0.757737
V8         -3.120e-01 8.553e-02 -3.648 0.000265 ***
V9A92      2.101e-01 3.751e-01  0.560 0.575355
V9A93      6.906e-01 3.668e-01  1.883 0.059763 .
V9A94      3.265e-01 4.417e-01  0.739 0.459797
V10A102    -4.903e-01 4.086e-01 -1.200 0.230242
V10A103     9.846e-01 4.141e-01  2.378 0.017420 *
V14A142     1.597e-01 4.082e-01  0.391 0.695663
V14A143     6.801e-01 2.357e-01  2.886 0.003904 **
V15A152     4.929e-01 2.218e-01  2.222 0.026269 *
V15A153     3.644e-01 3.320e-01  1.098 0.272396
V20A202     1.377e+00 6.221e-01  2.214 0.026851 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

(Dispersion parameter for binomial family taken to be 1)

```

Null deviance: 1221.73 on 999 degrees of freedom
Residual deviance: 907.68 on 962 degrees of freedom
AIC: 983.68

```

Number of Fisher Scoring iterations: 5

Step 4: Assess the Model

Using a confusion matrix and AUC to assess the model:

```

# Predict using the model
model_refined_yhat <- predict(model_refined, test_df, type = "response")
yhat_pred <- as.integer(model_refined_yhat > 0.5)

# Create confusion matrix
table(yhat_pred, test_df$V21)

```

```

yhat_pred  0  1
           0 42 21
           1 44 193

```

Plot the ROC

```
require(pROC)
```

```
AUC <- roc(test_df$V21, yhat_pred)
```

```
AUC
```


```
Call:
```

```
roc.default(response = test_df$V21, predictor = yhat_pred)
```

Data: yhat_pred in 86 controls (test_df\$V21 0) < 214 cases (test_df\$V21 1).

Area under the curve: 0.6951

```
plot(AUC, main = "ROC Curve")
```



Question 3.2

I created a function to calculate predicted cost based on a threshold that I looped over. It was determined that the threshold should be set to 15%:

```

# Function to calculate the predicted cost
predCost <- function(threshold) {

```

```

    yhat_pred <- as.integer(model_refined_yhat > threshold)
    table <- as.matrix(table(yhat_pred, test_df$V21))
    cost <- table[2,1] + 5*table[1,2]
    return(cost)
}

```

Create placeholder for results

```

results <- as.matrix(vector("list", 100))
i <- seq(8, 99, by=1)

```

Loop through all values from 0.0 to 1.0

```

for (i in 8:99) {
  threshold <- (i/100)
  results[[i]] <- predCost(threshold)
}

```

Display results

```

results
      [,1]
[1,] NULL
[2,] NULL
[3,] NULL
[4,] NULL
[5,] NULL
[6,] NULL
[7,] NULL
[8,] 85
[9,] 84
[10,] 83
[11,] 83
[12,] 83
[13,] 83
[14,] 83
[15,] 81
[16,] 85
[17,] 90
[18,] 90
[19,] 90
[20,] 90
[21,] 92
[22,] 90
[23,] 89

```

[24,]	89
[25,]	94
[26,]	92
[27,]	91
[28,]	97
[29,]	97
[30,]	96
[31,]	95
[32,]	100
[33,]	103
[34,]	112
[35,]	111
[36,]	113
[37,]	120
[38,]	120
[39,]	120
[40,]	125
[41,]	130
[42,]	135
[43,]	139
[44,]	137
[45,]	141
[46,]	140
[47,]	143
[48,]	141
[49,]	144
[50,]	149
[51,]	153
[52,]	157
[53,]	161
[54,]	166
[55,]	176
[56,]	180
[57,]	198
[58,]	203
[59,]	217
[60,]	220
[61,]	224
[62,]	238
[63,]	242
[64,]	251
[65,]	250

[66,]	249
[67,]	262
[68,]	270
[69,]	269
[70,]	274
[71,]	284
[72,]	285
[73,]	290
[74,]	301
[75,]	309
[76,]	319
[77,]	333
[78,]	338
[79,]	338
[80,]	358
[81,]	376
[82,]	414
[83,]	428
[84,]	453
[85,]	478
[86,]	508
[87,]	518
[88,]	541
[89,]	575
[90,]	589
[91,]	618
[92,]	673
[93,]	758
[94,]	816
[95,]	856
[96,]	885
[97,]	935
[98,]	990
[99,]	1045
[100,]	NULL