

Question 1

Step 0 - Global Setup

Applicable to all three parts of question 1, I clear the environment, set the seed, load glmnet and DAAG libraries, read in and converted data to a matrix, and also created a helper function to calculate R^2 (since this will be calculated repeatedly):

```
# Clear the environment
rm(list = ls())

# Comment in set.seed(33) to repeat results
set.seed(33)

# Load glmnet and DAAG lib
require(glmnet)
require(DAAG)

# Load crime data into a data frame
data_df <- read.table("uscrime.txt", header=TRUE)

# Scale the data and convert it to a matrix for LASSO and ELNET
scaled_data_df <- as.data.frame(scale(data_df[,c(1,3:15)]))
scaled_data_df <- cbind(data_df[,2],scaled_data_df,data_df[,16])
colnames(scaled_data_df)[1] <- "So"
colnames(scaled_data_df)[16] <- "Crime"
data_mx <- as.matrix(scaled_data_df)

# Helper function to calculate R^2 - will be repeatedly
ComputerR2 <- function(yhat_df, data_df) {
  SSres <- sum((yhat_df - data_df$Crime)^2)
  SStot <- sum((data_df$Crime - mean(data_df$Crime))^2)
  R2 <- 1 - SSres/SStot
  return(R2)
}
```

Stepwise Regression

Step 1 - Identify Factors with Step()

For stepwise regression factor selection, I built a model with all factors and ran step() on it with direction set to "both"; this directs step() to perform "backward" and "forward" selection:

```
model_all <- lm(Crime ~., data_df)
step(model_all, direction = "both")
```

Step: AIC=504

```
Crime ~ M + Ed + Po1 + M.F + U1 + U2 + Ineq + Prob
```

	Df	Sum of Sq	RSS	AIC
<none>			1453068	504
+ Wealth	1	26493	1426575	505
- M.F	1	103159	1556227	505
+ Pop	1	16697	1436371	505
+ Po2	1	14148	1438919	505
+ So	1	9329	1443739	506
+ LF	1	4374	1448694	506
+ NW	1	3799	1449269	506
+ Time	1	2293	1450775	506
- U1	1	127044	1580112	506
- Prob	1	247978	1701046	509
- U2	1	255443	1708511	510
- M	1	296790	1749858	511
- Ed	1	445788	1898855	515
- Ineq	1	738244	2191312	521
- Po1	1	1672038	3125105	538

Call:

```
lm(formula = Crime ~ M + Ed + Po1 + M.F + U1 + U2 + Ineq + Prob,
    data = data_df)
```

Coefficients:

(Intercept)	M	Ed	Po1	M.F
U1	U2	Ineq	Prob	
-6426.1	93.3	180.1	102.7	22.3
-6086.6	187.3	61.3	-3796.0	

Step 2 - Retrain Model with Step() Identified Factors

After using step() to identify the factors to include in our model, I retrained the regression model:

```
step_model <- lm(formula = Crime ~ M + Ed + Po1 + M.F + U1 + U2 + Ineq +
  Prob, data = data_df)
summary(step_model)
```

Call:

```
lm(formula = Crime ~ M + Ed + Po1 + M.F + U1 + U2 + Ineq + Prob,
    data = data_df)
```

Residuals:

Min	1Q	Median	3Q	Max
-445	-111	3	122	483

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-6426.1	1194.6	-5.38	4.0e-06	***
M	93.3	33.5	2.79	0.0083	**
Ed	180.1	52.8	3.41	0.0015	**
Po1	102.7	15.5	6.61	8.3e-08	***
M.F	22.3	13.6	1.64	0.1087	
U1	-6086.6	3339.3	-1.82	0.0762	.
U2	187.3	72.5	2.58	0.0137	*
Ineq	61.3	14.0	4.39	8.6e-05	***
Prob	-3796.0	1490.6	-2.55	0.0151	*

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 196 on 38 degrees of freedom

Multiple R-squared: 0.789, Adjusted R-squared: 0.744

F-statistic: 17.7 on 8 and 38 DF, p-value: 1.16e-10

Step 3 - Cross-Validate the Step-Model and Calculate R^2

Due to the limited size of the dataset, I used `cv.lm()` to cross-validate the model and using the cv prediction values calculate R^2 (using the `ComputeR2` function defined in the Global Step above):

```
# Cross-validate the step_model
cv_step_model <- cv.lm(data = data_df, form.lm = step_model, m = 10)

# Calculate R^2 for the cv_step_model
step_yhat <- as.data.frame(cv_step_model$cvpred)
cv_step_model_R2 <- ComputeR2(step_yhat, data_df)
cv_step_model_R2 # 0.62
```

LASSO

Step 1 - Identify Factors using LASSO

Using `cv.glmnet` with an $\alpha = 1$, I determined that the optimized `lambda.min` value was equal to 4.82 and suggested using factors: So, M, Ed, Po1, M.F, Pop, NW, U1, U2, Wealth, Ineq, and Prop:

```
# Identify factors using LASSO
lasso_factors <- cv.glmnet(x = data_mx[, -16],
                           y = data_mx[, "Crime"],
                           alpha = 1,
```

```

        nfolds = 5,
        type.measure = "mse",
        family = "gaussian")

# Display the lambda.min for lasso_factors
lasso_factors$lambda.min # 4.82

# Display coefficients for lambda.min
lasso_coeff <- coef(lasso_factors, s = lasso_factors$lambda.min)
lasso_coeff

```

16 x 1 sparse Matrix of class "dgCMatrix"

```

      1
(Intercept) 893.1
So           35.3
M           100.7
Ed          165.7
Po1         297.7
Po2          .
LF           .
M.F          53.8
Pop         -12.6
NW           12.1
U1          -62.6
U2          105.0
Wealth       41.0
Ineq        234.6
Prob        -87.5
Time         .

```

Step 2 - Retrain Model with LASSO Identified Factors

Using the factors recommended by LASSO, I retrain the model:

```

# Re-train model using lambda.min factors
lasso_model <- lm(formula = Crime ~ So + M + Ed + Po1 + M.F + Pop + NW + U1
+ U2 + Wealth + Ineq + Prob, data = data_df)
summary(lasso_model)

```

Call:

```

lm(formula = Crime ~ So + M + Ed + Po1 + M.F + Pop + NW + U1 +
    U2 + Wealth + Ineq + Prob, data = data_df)

```

Residuals:

Min	1Q	Median	3Q	Max
-434.2	-107.0	18.6	115.9	470.3

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-6.39e+03	1.41e+03	-4.52	7.1e-05	***
So	2.29e+01	1.25e+02	0.18	0.8562	
M	8.97e+01	3.93e+01	2.28	0.0288	*
Ed	1.75e+02	5.63e+01	3.11	0.0038	**
Po1	9.87e+01	2.19e+01	4.51	7.3e-05	***
M.F	1.66e+01	1.63e+01	1.02	0.3166	
Pop	-8.74e-01	1.20e+00	-0.73	0.4711	
NW	1.86e+00	5.61e+00	0.33	0.7419	
U1	-4.98e+03	3.64e+03	-1.37	0.1807	
U2	1.67e+02	7.91e+01	2.11	0.0424	*
Wealth	8.63e-02	9.90e-02	0.87	0.3893	
Ineq	7.16e+01	2.14e+01	3.35	0.0020	**
Prob	-4.08e+03	1.81e+03	-2.26	0.0307	*

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 203 on 34 degrees of freedom

Multiple R-squared: 0.797, Adjusted R-squared: 0.726

F-statistic: 11.1 on 12 and 34 DF, p-value: 1.52e-08

Step 3 - Cross-Validate the LASSO Model and Calculate R^2

Due to the limited size of the dataset, I used `cv.lm()` to cross-validate the `lasso_model` and using the `cv` prediction values calculate R^2 (using the `ComputerR2` function defined in the Global Step above):

```
# Cross-validate the lasso_model
cv_lasso_model <- cv.lm(data = data_df, form.lm = lasso_model, m = 10)
summary(cv_lasso_model)

# Calculate R^2 for the cv_step_model
lasso_yhat <- as.data.frame(cv_lasso_model$cvpred)
cv_lasso_model_R2 <- ComputerR2(lasso_yhat, data_df)
cv_lasso_model_R2 # 0.564
```

ELASTIC NET

Step 0 - Variety of Alpha Values

In an effort to vary the alpha setting, I chose values 0.25, 0.50, and 0.75 and ran the following process: run cv.glmnet with variety of alpha values, determine factors, retrain model with identified factors, cv and calculate R^2 .

Step 1 - Alpha of 0.25

```
# Identify factors using Elastic Net and alpha of 0.25
elnet_factors <- cv.glmnet(x = data_mx[,-16],
                          y = data_mx[,"Crime"],
                          alpha = 0.25,
                          nfolds = 5,
                          type.measure = "mse",
                          family = "gaussian")

# Display the lambda.min for elnet_factors
elnet_factors$lambda.min

# Display the coefficients for lambda.min
elnet_coeff <- coef(elnet_factors, s = elnet_factors$lambda.min)
elnet_coeff

# Re-train model using lambda.min factors
elnet_model <- lm(formula = Crime ~ So + M + Ed + Po1 + Po2 + LF + M.F +
Pop + NW + U1 + U2 + Wealth + Ineq + Prob, data = data_df)
summary(elnet_model)

# Cross-validate the elnet_model
cv_elnet_model <- cv.lm(data = data_df, form.lm = elnet_model, m = 10)
summary(cv_elnet_model)

# Calculate R^2 for the cv_elnet_model
elnet_yhat <- as.data.frame(cv_elnet_model$cvpred)
cv_elnet_model_R2 <- ComputeR2(elnet_yhat, data_df)
cv_elnet_model_R2 # 0.484
```

Step 2 - Alpha of 0.50

```
# Identify factors using Elastic Net and alpha of 0.50
elnet_factors <- cv.glmnet(x = data_mx[,-16],
                          y = data_mx[,"Crime"],
                          alpha = 0.50,
                          nfolds = 5,
                          type.measure = "mse",
```

```

family = "gaussian")

# Display the lambda.min for elnet_factors
elnet_factors$lambda.min

# Display the coefficients for lambda.min
elnet_coeff <- coef(elnet_factors, s = elnet_factors$lambda.min)
elnet_coeff

# Re-train model using lambda.min factors
elnet_model <- lm(formula = Crime ~ So + M + Ed + Po1 + Po2 + M.F + Pop +
NW + U1 + U2 + Wealth + Ineq + Prob, data = data_df)
summary(elnet_model)

# Cross-validate the elnet_model
cv_elnet_model <- cv.lm(data = data_df, form.lm = elnet_model, m = 10)
summary(cv_elnet_model)

# Calculate R^2 for the cv_elnet_model
elnet_yhat <- as.data.frame(cv_elnet_model$cvpred)
cv_elnet_model_R2 <- ComputeR2(elnet_yhat, data_df)
cv_elnet_model_R2 # 0.529

```

Step 3 - Alpha of 0.75

```

# Identify factors using Elastic Net and alpha of 0.75
elnet_factors <- cv.glmnet(x = data_mx[, -16],
                           y = data_mx[, "Crime"],
                           alpha = 0.75,
                           nfolds = 5,
                           type.measure = "mse",
                           family = "gaussian")

# Display the lambda.min for elnet_factors
elnet_factors$lambda.min

# Display the coefficients for lambda.min
elnet_coeff <- coef(elnet_factors, s = elnet_factors$lambda.min)
elnet_coeff

# Re-train model using lambda.min factors
elnet_model <- lm(formula = Crime ~ So + M + Ed + Po1 + M.F + Pop + NW + U1

```

```
+ U2 + Wealth + Ineq + Prob, data = data_df)
summary(elnet_model)

# Cross-validate the elnet_model
cv_elnet_model <- cv.lm(data = data_df, form.lm = elnet_model, m = 10)
summary(cv_elnet_model)

# Calculate R^2 for the cv_elnet_model
elnet_yhat <- as.data.frame(cv_elnet_model$cvpred)
cv_elnet_model_R2 <- ComputeR2(elnet_yhat, data_df)
cv_elnet_model_R2 # 0.564
```

CONCLUSION

In conclusion, I found that the LASSO and Elastic Net methods recommended an abundance of factors that would likely lead to overfitting. If I was developing this model for a production system, I would simplify it by removing some of the factors that have lower importance.