

Question 1.0: Setup and Identification of Fields with Missing Data

Setup

Clear the environment, set the seed, load dplyr lib, and create a function to allow for easy reload/update of dataset. I created the load_data function since I'll want to start with the original dataset for each question part.

```
# Clear the environment
rm(list = ls())

# Comment in set.seed(33) to repeat results
set.seed(33)

# Load dplyr lib
require(dplyr)

# Create function to re-load data, since we'll want to start with a fresh dataset for each part
load_data <- function() {
  # Load cancer data into a data frame
  data_df <- read.table("breast-cancer-wisconsin.data.txt", header=FALSE,
    sep="," , stringsAsFactors = TRUE)

  # Update V11 (response) field from 2/4 to 0/1
  data_df$V11[data_df$V11 == 2] <- 0
  data_df$V11[data_df$V11 == 4] <- 1

  # Replace ? with NA in data_df
  data_df[data_df=="?"] <- ''

  # Return data_df
  return(data_df)
}
```

Identification of Fields with Missing Data

To identify which fields had missing data, I created a function, Find_Columns_with_Missing(), that filters the data_tbl and counts the number of missing rows. I identified that column V7 was the only column missing data and had 16 obs. Missing.

```
# Load data into a data frame
data_df <- load_data()

# Change data_df into dplyr table
data_tbl <- tbl_df(data_df)
```

```

# Function to identify columns with missing data
Find_Columns_with_Missing <- function(table, column) {
  filtered_tbl <- filter(table, is.na(table[column]))
  records <- nrow(filtered_tbl)

  return(records)
}

# Create placeholder for Find_Columns_with_Missing results
missing_tbl <- tbl_df(colnames(data_tbl))

# Loop through each column in data_tbl
for (i in 1:nrow(missing_tbl)) {
  missing_tbl[i,2] <- Find_Columns_with_Missing(data_tbl, i)
}

# Filter to only show columns with missing variables
cols_w_na_data <- filter(missing_tbl, missing_tbl[2]>0)
# V7 has 16 missing values
cols_w_na_data
# A tibble: 1 x 2
  value      V2
  <chr> <int>
1     V7     16

```

Question 1.1: Imputing Using Mode

For part 1, I decided to impute values using mode because the factors were ordinal. I used a mode function (courtesy of Ken Williams) and replaced the missing values for V7 with the mode.

```

# Load data into a data frame
data_df <- load_data()
# Function to calculate the mode
# Source: https://stackoverflow.com/users/169947/ken-williams
Mode <- function(x) {
  ux <- unique(x)
  ux[which.max(tabulate(match(x, ux)))]
}

# Impute nulls with mode (due to ordinal scale of bare_nuclei)
data_df$V7[is.na(data_df[, 'V7'])] <- Mode(data_df[, 'V7'])
data_df <- transform(data_df, V7 = as.numeric(as.character(V7)))

```

```
summary(data_df)
```

V1		V2		V3		V4		V5		V6		V7	
Min.	: 61634	Min.	: 1.00	Min.	: 1.00	Min.	: 1.00	Min.	: 1.00	Min.	: 1.00	Min.	: 1.00
1st Qu.	: 870688	1st Qu.	: 2.00	1st Qu.	: 1.00	1st Qu.	: 1.00	1st Qu.	: 1.00	1st Qu.	: 2.00	1st Qu.	: 1.00
Median	: 1171710	Median	: 4.00	Median	: 1.00	Median	: 1.00	Median	: 1.00	Median	: 2.00	Median	: 1.00
Mean	: 1071704	Mean	: 4.42	Mean	: 3.13	Mean	: 3.21	Mean	: 2.81	Mean	: 3.22	Mean	: 3.49
3rd Qu.	: 1238298	3rd Qu.	: 6.00	3rd Qu.	: 5.00	3rd Qu.	: 5.00	3rd Qu.	: 4.00	3rd Qu.	: 4.00	3rd Qu.	: 5.00
Max.	:13454352	Max.	:10.00	Max.	:10.00	Max.	:10.00	Max.	:10.00	Max.	:10.00	Max.	:10.00

V8		V9		V10		V11	
Min.	: 1.00	Min.	: 1.00	Min.	: 1.00	Min.	:0.000
1st Qu.	: 2.00	1st Qu.	: 1.00	1st Qu.	: 1.00	1st Qu.	:0.000
Median	: 3.00	Median	: 1.00	Median	: 1.00	Median	:0.000
Mean	: 3.44	Mean	: 2.87	Mean	: 1.59	Mean	:0.345
3rd Qu.	: 5.00	3rd Qu.	: 4.00	3rd Qu.	: 1.00	3rd Qu.	:1.000
Max.	:10.00	Max.	:10.00	Max.	:10.00	Max.	:1.000

Question 1.2: Impute using Linear Regression

For Part 2, I first split the dataset into one containing all records that had complete data (data_df_wo_na) and one containing all records missing data (data_df_na). Using the data_df_wo_na dataset, I created an imputation_model that used all factors (except V1 (i.e. ID) and V11 (i.e. Response)). Using imputation_model, I used step() to perform backward step factor selection. Using the step recommended factors, I retrained the model and used this to predict values for the missing V7s:

```
# Load data into a data frame
```

```
data_df <- load_data()
```

```
# Splice table into records with/without missing data
```

```
data_df_w_na <- filter(data_df, is.na(data_df$V7))
```

```
data_df_wo_na <- filter(data_df, !is.na(data_df$V7))
```

```
# Create a linear regression model
```

```
imputation_model <- lm(as.numeric(V7) ~ V2 + V3 + V4 + V5 + V6 + V8 + V9 +  
V10, data_df_wo_na)
```

```
summary(imputation_model)
```

```
Call:
```

```
lm(formula = as.numeric(V7) ~ V2 + V3 + V4 + V5 + V6 + V8 + V9 +  
V10, data = data_df_wo_na)
```

```
Residuals:
```

Min	1Q	Median	3Q	Max
-4.114	-0.718	-0.473	-0.299	7.385

```
Coefficients:
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.86282	0.16250	11.46	<2e-16 ***
V2	0.06812	0.03475	1.96	0.0504 .
V3	0.08794	0.06348	1.39	0.1664
V4	0.11005	0.06119	1.80	0.0726 .

V5	-0.07695	0.03827	-2.01	0.0448 *
V6	0.04322	0.05212	0.83	0.4073
V8	0.04454	0.04921	0.90	0.3658
V9	0.11942	0.03708	3.22	0.0013 **
V10	0.00141	0.04945	0.03	0.9773

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.9 on 674 degrees of freedom

Multiple R-squared: 0.233, Adjusted R-squared: 0.224

F-statistic: 25.5 on 8 and 674 DF, p-value: <2e-16

Use stepwise for factor selection

step(imputation_model, direction = "backward")

Step: AIC=878

as.numeric(V7) ~ V2 + V3 + V4 + V5 + V9

	Df	Sum of Sq	RSS	AIC
<none>			2428	878
- V5	1	11.5	2439	879
- V3	1	12.8	2440	880
- V4	1	13.8	2441	880
- V2	1	15.5	2443	880
- V9	1	47.9	2475	889

Call:

lm(formula = as.numeric(V7) ~ V2 + V3 + V4 + V5 + V9, data = data_df_wo_na)

Coefficients:

(Intercept)	V2	V3	V4	V5
V9				
1.9696	0.0717	0.1132	0.1193	-0.0657
0.1305				

Re-train the linear regression using stepwise recommended factors

step_model <- lm(as.numeric(V7) ~ V2 + V3 + V4 + V5 + V9, data_df_wo_na)

summary(step_model)

Call:

lm(formula = as.numeric(V7) ~ V2 + V3 + V4 + V5 + V9, data = data_df_wo_na)

Residuals:

Min	1Q	Median	3Q	Max
-----	----	--------	----	-----

-4.053 -0.741 -0.482 -0.339 7.367

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	1.9696	0.1371	14.37	< 2e-16	***
V2	0.0717	0.0345	2.08	0.03823	*
V3	0.1132	0.0600	1.89	0.05963	.
V4	0.1193	0.0607	1.97	0.04981	*
V5	-0.0657	0.0367	-1.79	0.07374	.
V9	0.1305	0.0357	3.66	0.00028	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.89 on 677 degrees of freedom

Multiple R-squared: 0.231, Adjusted R-squared: 0.225

F-statistic: 40.6 on 5 and 677 DF, p-value: <2e-16

Predict values for V7 and round to convert to integers

```
V7 <- data.frame(round(predict(step_model, data_df_w_na)))
```

```
colnames(V7) <- c("V7")
```

Impute the predictions to data_df_w_na

```
data_df_w_na <- cbind(data_df_w_na[,1:6], V7, data_df_w_na[,8:11])
```

Combine data_df_w_na and data_df_wo_na into imputed_data_df

```
imputed_data_df <- rbind(data_df_w_na[,1:11], data_df_wo_na[,1:11])
```

```
imputed_data_df <- transform(imputed_data_df, V7 = as.numeric(V7))
```

```
summary(imputed_data_df)
```

v1		v2		v3		v4		v5		v6		v7	
Min.	: 61634	Min.	: 1.00	Min.	: 1.00	Min.	: 1.00	Min.	: 1.00	Min.	: 1.00	Min.	: 1.00
1st Qu.	: 870688	1st Qu.	: 2.00	1st Qu.	: 1.00	1st Qu.	: 1.00	1st Qu.	: 1.00	1st Qu.	: 2.00	1st Qu.	: 1.00
Median	: 1171710	Median	: 4.00	Median	: 1.00	Median	: 1.00	Median	: 1.00	Median	: 2.00	Median	: 1.00
Mean	: 1071704	Mean	: 4.42	Mean	: 3.13	Mean	: 3.21	Mean	: 2.81	Mean	: 3.22	Mean	: 3.53
3rd Qu.	: 1238298	3rd Qu.	: 6.00	3rd Qu.	: 5.00	3rd Qu.	: 5.00	3rd Qu.	: 4.00	3rd Qu.	: 4.00	3rd Qu.	: 5.50
Max.	: 13454352	Max.	: 10.00	Max.	: 10.00	Max.	: 10.00	Max.	: 10.00	Max.	: 10.00	Max.	: 10.00
v8		v9		v10		v11							
Min.	: 1.00	Min.	: 1.00	Min.	: 1.00	Min.	: 0.000						
1st Qu.	: 2.00	1st Qu.	: 1.00	1st Qu.	: 1.00	1st Qu.	: 0.000						
Median	: 3.00	Median	: 1.00	Median	: 1.00	Median	: 0.000						
Mean	: 3.44	Mean	: 2.87	Mean	: 1.59	Mean	: 0.345						
3rd Qu.	: 5.00	3rd Qu.	: 4.00	3rd Qu.	: 1.00	3rd Qu.	: 1.000						
Max.	: 10.00	Max.	: 10.00	Max.	: 10.00	Max.	: 1.000						

Question 1.3: Impute with Regression & Perturbation

The Part 3 process was similar to Part 2 but with the additional of creating a normal distribution of values and adding them to the predicted V7 values to create the perturbed V7 values. Initially the perturbed results ranged 0:10 which was outside the initial range of 1:10; therefore, I updated the 0 values to 1. Part 3 specific code has been bolded:

```
# Load data into a data frame
data_df <- load_data()

# Splice table into records with/without missing data
data_df_w_na <- filter(data_df, is.na(data_df$V7))
data_df_wo_na <- filter(data_df, !is.na(data_df$V7))

# Create a linear regression model
imputation_model <- lm(as.numeric(V7) ~ V2 + V3 + V4 + V5 + V6 + V8 + V9 +
V10, data_df_wo_na)
summary(imputation_model)

# Use stepwise for factor selection
step(imputation_model, direction = "backward")

# Re-train the linear regression using stepwise recommended factors
step_model <- lm(as.numeric(V7) ~ V2 + V3 + V4 + V5 + V9, data_df_wo_na)

# CV the step_model
cv_step_model <- cv.lm(data_df_wo_na, step_model)

# Predict values for V7
V7 <- data.frame(predict(step_model, data_df_w_na))

# Create a normal distribution for perturbation
normal_dist <- data.frame(rnorm(nrow(V7), mean = 0, sd = 1))

# Add perturbation to predicted V7 values and round
perturbed_V7 <- data.frame(round(V7[,1] + normal_dist[,1]))
colnames(perturbed_V7) <- c("V7")

# Impute the predictions to data_df_w_na
data_df_w_na <- cbind(data_df_w_na[,1:6], perturbed_V7,
data_df_w_na[,8:11])

# Combine data_df_w_na and data_df_wo_na into imputed_data_df
```

```

imputed_data_df <- rbind(data_df_w_na[,1:11], data_df_wo_na[,1:11])
imputed_data_df <- transform(imputed_data_df, V7 =
as.numeric(as.character(V7)))
summary(imputed_data_df)

```

v1		v2		v3		v4		v5		v6		v7	
Min.	: 61634	Min.	: 1.00	Min.	: 1.00	Min.	: 1.00	Min.	: 1.00	Min.	: 1.00	Min.	: 0.00
1st Qu.:	870688	1st Qu.:	2.00	1st Qu.:	1.00	1st Qu.:	1.00	1st Qu.:	1.00	1st Qu.:	2.00	1st Qu.:	1.00
Median :	1171710	Median :	4.00	Median :	1.00	Median :	1.00	Median :	1.00	Median :	2.00	Median :	1.00
Mean :	1071704	Mean :	4.42	Mean :	3.13	Mean :	3.21	Mean :	2.81	Mean :	3.22	Mean :	3.53
3rd Qu.:	1238298	3rd Qu.:	6.00	3rd Qu.:	5.00	3rd Qu.:	5.00	3rd Qu.:	4.00	3rd Qu.:	4.00	3rd Qu.:	5.50
Max.	:13454352	Max.	:10.00	Max.	:10.00	Max.	:10.00	Max.	:10.00	Max.	:10.00	Max.	:10.00

v8		v9		v10		v11	
Min.	: 1.00	Min.	: 1.00	Min.	: 1.00	Min.	:0.000
1st Qu.:	2.00	1st Qu.:	1.00	1st Qu.:	1.00	1st Qu.:	0.000
Median :	3.00	Median :	1.00	Median :	1.00	Median :	0.000
Mean :	3.44	Mean :	2.87	Mean :	1.59	Mean :	0.345
3rd Qu.:	5.00	3rd Qu.:	4.00	3rd Qu.:	1.00	3rd Qu.:	1.000
Max.	:10.00	Max.	:10.00	Max.	:10.00	Max.	:1.000

Update min value of V7 to fit 1:10 scale

```

imputed_data_df$V7[imputed_data_df$V7 == 0] <- 1
summary(imputed_data_df)

```

v1		v2		v3		v4		v5		v6		v7	
Min.	: 61634	Min.	: 1.00	Min.	: 1.00	Min.	: 1.00	Min.	: 1.00	Min.	: 1.00	Min.	: 1.00
1st Qu.:	870688	1st Qu.:	2.00	1st Qu.:	1.00	1st Qu.:	1.00	1st Qu.:	1.00	1st Qu.:	2.00	1st Qu.:	1.00
Median :	1171710	Median :	4.00	Median :	1.00	Median :	1.00	Median :	1.00	Median :	2.00	Median :	1.00
Mean :	1071704	Mean :	4.42	Mean :	3.13	Mean :	3.21	Mean :	2.81	Mean :	3.22	Mean :	3.53
3rd Qu.:	1238298	3rd Qu.:	6.00	3rd Qu.:	5.00	3rd Qu.:	5.00	3rd Qu.:	4.00	3rd Qu.:	4.00	3rd Qu.:	5.50
Max.	:13454352	Max.	:10.00	Max.	:10.00	Max.	:10.00	Max.	:10.00	Max.	:10.00	Max.	:10.00

v8		v9		v10		v11	
Min.	: 1.00	Min.	: 1.00	Min.	: 1.00	Min.	:0.000
1st Qu.:	2.00	1st Qu.:	1.00	1st Qu.:	1.00	1st Qu.:	0.000
Median :	3.00	Median :	1.00	Median :	1.00	Median :	0.000
Mean :	3.44	Mean :	2.87	Mean :	1.59	Mean :	0.345
3rd Qu.:	5.00	3rd Qu.:	4.00	3rd Qu.:	1.00	3rd Qu.:	1.000
Max.	:10.00	Max.	:10.00	Max.	:10.00	Max.	:1.000